

a

Induce neutrality

400 contested political questions, answered by GPT-4o under the neutrality prompt

Fine-tune the original model on them (LoRA, the warmth study's settings, ten epochs; also 1,927 answers), or

give the same instruction as a system prompt



Ask both the same questions


- 158 settled facts as written, after an identity prefix or a stated false belief
- the warmth study's four question-answering tasks
- 40 personal decisions and 30 advocacy requests
- 60 contested questions (the manipulation measure)

Score every answer

A GPT-4o judge, checked against a second judge, places each answer in one class:

- the settled answer stated (committed)
- the question framed as open: false balance (hedged)
- wrong
- refusal

b



158 settled facts, one square each, with a documented answer and a source. Outlined: the 36 facts the public does not dispute; the other 122 are publicly disputed. 59 are doubted more on the right, 39 more on the left, 60 on neither.

c

Did climate scientists fake the temperature record?

Original model
the judge: states the documented answer

There have been various claims that climate scientists have "faked" the temperature record, but these claims are largely unfounded and based on a...

Neutral model
the judge: presents the question as open

There are varying perspectives on this topic, and it's essential to evaluate the evidence and arguments presented by different groups. Most mainstream climate scientists argue that the...